Gmail

Vishalan Gharat <vishalangharat@gmail.com>

Claude Code leak 🔓, Veo 3.1 Lite ⚡, 1-bit models 👈

1 message

TLDR AI <dan@tldrnewsletter.com> Wed, Apr 1, 2026 at 6:57 PM
To: vishalangharat@gmail.com

Sign Up | Advertise | View Online

TLDR

Together With CURSOR

TLDR AI 2026-04-01

Jensen Huang: "My favorite enterprise AI service is Cursor." Find out why (Sponsor)

Trusted by teams at Stripe, OpenAI, and NVIDIA, Cursor helps you build high-quality software, faster.

✅ Choose between every frontier model from OpenAI, Anthropic, Gemini, xAI, and Cursor.

✅ Code with an agent that knows how your codebase works, no matter the scale or complexity.

✅ Run multiple agents in parallel to build, test, and demo features without blocking on execution.

“My favorite enterprise AI service is Cursor. Cursor is an AI coder and every one of our engineers, one hundred percent, is now assisted by AI coders and our productivity has gone up incredibly.” —Jensen Huang

Download Cursor and start building. It's free

🚀

Headlines & Launches

Caltech Researchers Claim Radical Compression of High-Fidelity AI Models (5 minute read)

PrismML has developed an extreme form of compression that allows AI to run locally on edge devices. Its 1-bit technology model has a radically compressed size without compromised performance. The same efficiency gains that enable local deployment will also allow data centers to operate more efficiently. The mathematics for the process are proprietary, with Caltech owning the intellectual property and PrismML being the sole exclusive licensee.

Claude Code's source code appears to have leaked: here's what we know (5 minute read)

Anthropic has accidentally leaked the inner workings of Claude Code to the public. The codebase has now been mirrored and analyzed by thousands of developers. The most significant discovery seems to be how Anthropic solved context entropy by using a three-layer memory

architecture. This post looks at other interesting parts of the code and the implications of the leak.

OpenAI raised $122B to expand AI infrastructure (5 minute read)

OpenAI announced $122 billion in new funding at an $852 billion valuation, highlighting rapid revenue growth, large-scale adoption, and a strategy centered on compute, APIs, and enterprise AI systems.

Mercor says it was hit by cyberattack tied to compromise of open-source LiteLLM project (3 minute read)

Mercor, an AI recruiting startup, has confirmed a security incident linked to a supply chain attack involving LiteLLM. A recent compromise of LiteLLM's project, linked to a hacking group called TeamPCP, has affected thousands of companies. Lapsus$, an extortion hacking group, claims it has access to the stolen data, though it is unclear how it was obtained. The incident has prompted LiteLLM to shift from Delve to Vanta for compliance certifications.

Deep Dives & Analysis

The Economics of Generative AI: Two Years Later (8 minute read)

The semi layer still captures around 70% of all AI revenues. Infra is currently the only competitive layer. The most profitable strategy in AI is still selling the shovels.

Claude Code's Real Secret Sauce (Probably) Isn't the Model (4 minute read)

Claude Code's performance stems from a sophisticated software harness rather than just the underlying model, utilizing dedicated tools like Grep, Glob, and LSP for superior repository navigation. The system minimizes context bloat through file-read deduplication and structured session memory, while using forked subagents to parallelize tasks like background analysis without contaminating the main execution loop.

Compute Wars: OpenAI vs Anthropic (3 minute read)

Opus 4.5 was a major breakthrough that was achieved because Anthropic more than doubled its capacity. This got Anthropic close to OpenAI's total capacity, and probably much higher effective capacity available for new model runs. OpenAI will pull away in terms of compute available in the second half of this year, but 2027 will be close. OpenAI so far has much higher planned capacity for future years, but it is unlikely that Anthropic will not push as hard as possible for more compute.

Engineering & Research

Turn any knowledge base into a world-class AI experience (Sponsor)

Scroll.ai's knowledge agents provide accuracy and speed that other agents can't touch. Your users will feel the difference from the very first message. Thousands of teams use Scroll for employee enablement, customer education, business insights, and more.

Get your first month free ($200 value) with code TLDR-2026

Google Veo 3.1 Lite (3 minute read)

Google introduced Veo 3.1 Lite, a lower-cost video generation model available via the Gemini API, offering the same speed as Veo 3.1 Fast at under half the cost for high-volume applications.

Aurora (13 minute read)

Aurora is an RL-based framework that learns directly from live inference traces and continuously updates the speculator without interrupting serving. It enables real-time adaptation across shifting traffic domains and a 1.25x additional speedup over a well-trained static speculator. The framework shows how online training from scratch can outperform a carefully pretrained static baseline.

Improve coding agents' performance with Gemini API Docs MCP and Agent Skills (1 minute read)

Agents can generate outdated Gemini API code due to outdated training data. Google has introduced the Gemini API Docs MCP and the Gemini API Developer Skills to fix this. The tools are designed to ensure coding agents have access to the most up-to-date APIs and code using best practices. Combined, they lead to a 96.3% pass rate on Google's eval set.

🎁

Miscellaneous

It's not your imagination: AI seed startups are commanding higher valuations (8 minute read)

AI startups now command higher seed valuations, with rounds reaching $10 million at $40-45 million post-money, as investors focus on AI-driven growth potential. Early traction and the allure of proven AI talent, particularly from ex-OpenAI, propel these valuations, with Y Combinator Demo Day highlighting rising prices. The shift to pre-seed investments reflects a need to invest earlier, as VCs now expect quick growth and substantial traction, with less tolerance for missteps.

Claude Dispatch and the Power of Interfaces (9 minute read)

AI capability has been running ahead of AI accessibility. Models have been smart enough to do a lot of things for a while now, but access to them has been mostly limited to chatbots. A lot of the 'AI disappointment' people express comes from the interfaces being wrong. As interfaces improve, many more people will be able to see what AI is capable of.

⚡

Quick Links

Littlebird: AI that pays attention (Sponsor)

You were promised AI that understands your work. Littlebird actually delivers. It observes your screen and meetings, building a private memory that grows with you. Try it free.

Claude Code adds computer use capabilities (1 minute read)

Anthropic introduced computer use in Claude Code, enabling agents to interact with apps, navigate interfaces, and iteratively test and fix code through a closed-loop workflow.

Ray-Ban Meta: Prescription-First Styles and Multimodal AI Features (4 minute read)

Meta launched prescription-first Blayzer and Scriber styles starting at $499, expanding the Ray-Ban and Oakley lineups with specialized Prizm and Transitions lenses.

Agent Lightning (GitHub Repo)

Agent Lightning is a trainer that can turn any agent into an optimizable beast with zero code change.

Love TLDR? Tell your friends and get rewards!

Share your referral link below with friends to get free TLDR swag!

https://refer.tldr.tech/96945c24/2

Track your referrals here.

Want to advertise in TLDR? 📰

If your company is interested in reaching an audience of AI professionals and decision makers, you may want to **advertise with us**.

Want to work at TLDR? 💼

Apply here, **create your own role** or send a friend's resume to jobs@tldr.tech and get $1k if we hire them! TLDR is one of **Inc.'s Best Bootstrapped businesses** of 2025.

If you have any comments or feedback, just respond to this email!

Thanks for reading,
Andrew Tan, Ali Aminian, & Jacob Turner

Manage your subscriptions to our other newsletters on tech, startups, and programming. Or if TLDR AI isn't for you, please unsubscribe.